

CS747 : Programming Assignment 2

Adityaya Dhande 210070005

October 13, 2023

Task 1

UCB




Figure 1: Regret vs Horizon for UCB algorithm

Regret increases with horizon as expected and follows a logarithmic variation

```
1 class UCB(Algorithm):
2     def __init__(self, num_arms, horizon):
3         super().__init__(num_arms, horizon)
4         self.ucbs = np.zeros(num_arms)
5         self.u = np.zeros(num_arms)
6         self.p_hat = np.zeros(num_arms)
7
8     def give_pull(self):
9         t = sum(self.u)
```

```

10     self.ucbs = self.p_hat + np.sqrt(2 * np.log(t) / self.u)
11     return np.argmax(self.ucbs)
12
13     def get_reward(self, arm_index, reward):
14         self.u[arm_index] += 1
15         n = self.u[arm_index]
16         mean = self.p_hat[arm_index]
17         new_mean = ((n - 1) / n) * mean + (reward / n)
18         self.p_hat[arm_index] = new_mean

```

Code explanation : `self.ucbs` is an array containing the UCBs of all the arms in a given iteration. `self.u` contains the number of times each arm has been pulled till the current iteration. `self.p_hat` contains the empirical means of all the arms based on pulls till the current iteration.

The `get_reward` function takes the index of the arm which was pulled and the reward and increments `self.u[arm_index]` by 1 because that arm was pulled in this iteration, and thus the total number of pulls of that arm has increased by 1. If now the total number of pulls is n then the old number of pulls was $n - 1$, and thus the total reward before the current pull was $\text{old mean} \times (n - 1)$. The new mean is

$$\text{new mean} = \frac{\text{total reward}}{\text{total number of pulls}} = \frac{\text{total reward before current pull} + \text{reward of current pull}}{\text{total number of pulls}}$$

The new mean is set for the arm that was pulled.

The `give_pull` function calculates the UCB for each arm as $ucb_a^t = \hat{p}_a^t + \sqrt{\frac{2 \ln(t)}{u_a^t}}$ and returns the index of the arm with the maximum UCB.

KL-UCB




Figure 2: Regret vs Horizon for KL-UCB

Regret increases with horizon as expected and follows a logarithmic variation. The decrease in the end can be explained by negative regret generated because of the optimal arm being pulled many times.

```

1  def KL(x ,y):
2  if x == 0 :
3      return math.log(1/(1 - y))
4  elif x == 1 :
5      return math.log(1/y)
6  return (x * math.log(x / y) + (1 - x) * math.log((1 - x) / (1 - y)))
7
8  def KL_ucb(p, u_a, t, c = 3, tol = 1e-3) :
9      l = p
10     u = 1
11     q = (l + u)/2
12     target = (math.log(t) + c * math.log(math.log(t)))/u_a
13     while (u - l > tol):
14         q = (l + u)/2
15         current = KL(p,q)
16         if current < target :
17             l = q
18         elif current > target :
19             u = q
20         else :
21             return q
22     return q
23
24 class KL_UCB(Algorithm):
25 def __init__(self, num_arms, horizon):
26     super().__init__(num_arms, horizon)
27     self.first_pull = True
28     self.kl_ucbs = np.zeros(num_arms)
29     self.u = np.zeros(num_arms)
30     self.p_hat = np.zeros(num_arms)
31
32 def give_pull(self):
33     if self.first_pull :
34         arm = int(sum(self.u))
35         if arm == (len(self.kl_ucbs) - 1) :
36             self.first_pull = False
37         return arm
38     t = sum(self.u)
39     for i in range(len(self.kl_ucbs)):
40         self.kl_ucbs[i] = KL_ucb(self.p_hat[i], self.u[i], t, c=0)
41     arm = np.argmax(self.kl_ucbs)
42     return arm
43
44 def get_reward(self, arm_index, reward):
45     self.u[arm_index] += 1
46     n = self.u[arm_index]
47     mean = self.p_hat[arm_index]
48     new_mean = ((n - 1) / n) * mean + (reward / n)
49     self.p_hat[arm_index] = new_mean

```

Code explanation : `self.first_pull` is a boolean variable which is `True` till all the arms have been sampled for the first time. `self.kl_ucbs` is an array containing the KL-UCBs of all the arms in a given iteration. `self.u` contains the number of times each arm has been pulled till the current iteration. `self.p_hat` contains the empirical means of all the arms based on pulls till the current iteration. The `get_reward` function takes the index of the arm which was pulled and the reward and increments `self.u[arm_index]` by 1 because that arm was pulled in this iteration, and thus the total number of pulls of that arm has increased by 1. If now the total number of pulls is n then the old number of pulls

was $n - 1$, and thus the total reward before the current pull was $\text{old mean} \times (n - 1)$. The new mean is

$$\text{new mean} = \frac{\text{total reward}}{\text{total number of pulls}} = \frac{\text{total reward before current pull} + \text{reward of current pull}}{\text{total number of pulls}}$$

The new mean is set for the arm that was pulled.

The `give_pull` function samples each arm once for the first `num_arms` times. From then on it calculates the KL-UCB for each arm as q such that

$$\text{self.u}[\text{arm_index}] \times KL(\text{self.p_hat}[\text{arm_index}], q) = \ln(t) + c \ln(\ln(t))$$

and returns the index of the arm with the highest KL-UCB. The value of c used here is 0. The value of q is found using binary search because $KL(p, p) = 0$ and $KL(p, 1) = \infty$ and it increases as the second argument increases from p to 1. The function `KL_ucb` does exactly this with a tolerance in the value of q of about 10^{-3} . The function `KL(x, y)` returns the KL divergence of 2 bernoulli distributions with means x and y .

Thompson Sampling

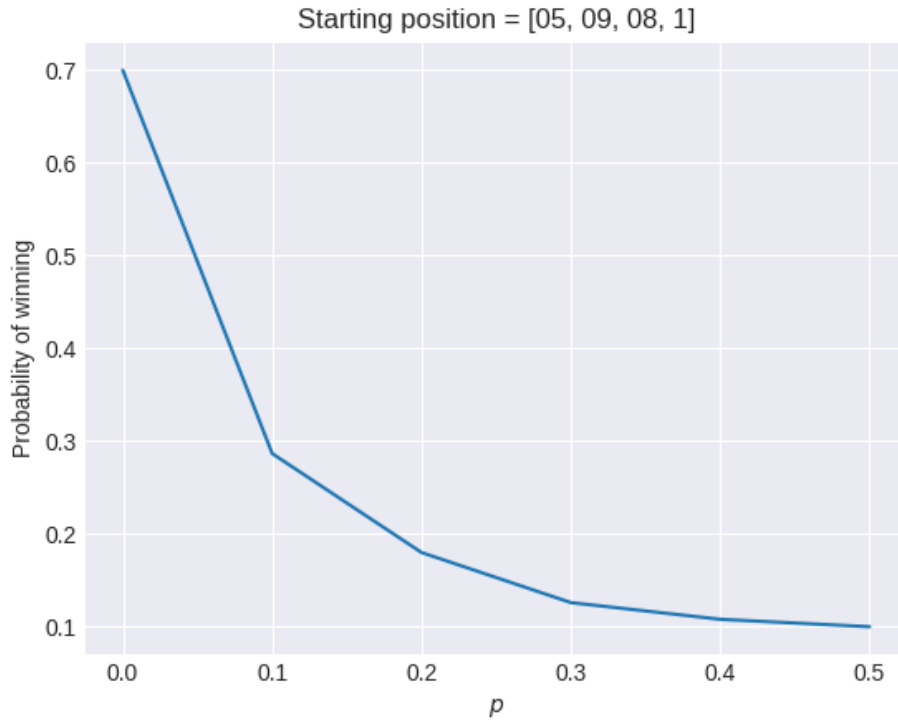


Figure 3: Regret vs Horizon for Thompson Sampling

Regret increases with horizon as expected

```

1 class Thompson_Sampling(Algorithm):
2     def __init__(self, num_arms, horizon):
3         super().__init__(num_arms, horizon)
4         self.sa = np.zeros(num_arms)
5         self.fa = np.zeros(num_arms)

```

```

6         self.t_samples = np.zeros(num_arms)
7
8     def give_pull(self):
9         for i in range(len(self.t_samples)) :
10             self.t_samples[i] = np.random.beta(self.sa[i] + 1, self.fa[i] + 1)
11         return np.argmax(self.t_samples)
12
13     def get_reward(self, arm_index, reward):
14         self.sa[arm_index] += reward
15         self.fa[arm_index] += 1 - reward

```

Code explanation : `self.sa` stores the number of successes(reward=1) of all the arms until the current iteration. `self.fa` stores the number of failures(reward=0) of all the arms until the current iteration. `self.t_samples` stores the samples drawn from the beta distribution corresponding to each arm.

The `get_reward` function takes the arm index and the reward and increases the successes of that arm by 1 if the reward is 1 and increases the failures of that arm by 1 if the reward is 0

The `give_pull` function draws, for each arm, samples from

$$\beta(\text{self.sa}[\text{arm_index}] + 1, \text{self.fa}[\text{arm_index}] + 1)$$

and returns the index of the arm whose β distribution gives the greatest sample.